

CS 355 Lecture 7 (4/23)

Logistics: HW2 due this Friday at 5pm (submit via Gradescope)

HW3 posted this Friday

See Piazza for office hour information for upcoming week (course staff is traveling)

Previous lecture: 2-party computation (Yao's garbled circuits)

Suppose Alice and Bob want to compute function f on joint input (x, y)

Alice (x)

OT requests for x

OT responses
garbled circuit for f
garbled inputs for y

$f(x, y)$

Bob (y)

requires algebraic assumptions
(more expensive)

only requires cheap symmetric
primitives

[HW2: Realize large number of OTs from
a small ($\sim \lambda$) number of base OTs]

Security: neither party learns more from protocol other than what is revealed by $f(x, y)$

$$\text{e.g. } \underbrace{\{\text{View}_A(x, y)\}}_{\text{view of Alice in protocol execution on inputs } x \text{ and } y} \stackrel{\epsilon}{\approx} \underbrace{\{S(1^\lambda, x, f(x, y))\}}_{\text{real protocol transcript can be simulated just given party's input and output of computation}}$$

view of Alice in
protocol execution
on inputs x and y

real protocol transcript
can be simulated just given
party's input and output of computation

Question: What if we have more than two parties?

This lecture: Secret sharing and MPC

Beaver triples and MPC in the preprocessing model

Secret sharing: Suppose we have a secret and want to distribute it among n parties such that any t of them can subsequently recover the secret and any $(t-1)$ subset cannot [e.g., Board of directors at Coca-Cola want to protect Coca-Cola recipe]

Def. A (t, n) -secret sharing scheme over a message space \mathcal{M} and share space \mathcal{S} consists of two efficient algorithms:

Share: $\mathcal{M} \rightarrow \mathcal{S}^n$

Reconstruct: $\mathcal{S}^t \rightarrow \mathcal{M}$

with the following properties:

Correctness: Any t shares can be used to reconstruct m :

$$\forall m \in \mathcal{M} : (s_1, \dots, s_n) \leftarrow \text{Share}(m)$$

$$\forall S \subseteq \{s_1, \dots, s_n\} \text{ where } |S| = t : \text{Reconstruct}(S) = m$$

Security: Need at least t shares to learn the secret

$$\forall m_0, m_1 \in \mathcal{M}, \forall I \subseteq [n] \text{ where } |I| < t :$$

$$\{(s_1, \dots, s_n) \leftarrow \text{Share}(m_0) : s_i \text{ for all } i \in I\} \equiv \{(s_1, \dots, s_n) \leftarrow \text{Share}(m_1) : s_i \text{ for all } i \in I\}$$

can relax to computational indistinguishability (in which case, Share takes in additional security parameter)

↪ notion here is information-theoretic

Examples: Additive secret sharing [n out of n]: $M = \mathbb{Z}_p$, $S = \mathbb{Z}_p^n$ any ring will work (p need not be prime)

Provides information-theoretic Security - check this!

- Share(m): Sample $r_1, \dots, r_{n-1} \xleftarrow{R} \mathbb{Z}_p$ and set $r_n = m = \sum_{i=1}^{n-1} r_i \in \mathbb{Z}_p$

- Reconstruct(r_1, \dots, r_n): Output $\sum_{i=1}^n r_i$

Combinatorial secret sharing [t out of n]: Will use a symmetric encryption scheme over $\{0,1\}^*$ (e.g. AES-CTR) satisfying CPA-security

- Share(m): Sample n keys k_1, \dots, k_n for encryption scheme

For every t-subset $\{i_1, \dots, i_t\} \subseteq [n]$, encrypt m using k_{i_1}, \dots, k_{i_t} (e.g. $\text{Enc}(k_{i_1}, \text{Enc}(k_{i_2}, \dots, \text{Enc}(k_{i_t}, m) \dots))$)

Let $\{ct\}$ be the collection of ciphertexts

Output shares $((k_1, \{ct\}), \dots, (k_n, \{ct\}))$

- Reconstruct $((k_1, \{ct\}), \dots, (k_t, \{ct\}))$: by construction, there is one $ct \in \{ct\}$ encrypted under k_1, \dots, k_t , so decrypt accordingly very large shares! Question: Can we do better?

relies on computational assumptions


Shamir secret sharing [t out of n]: $M = \mathbb{F}_p$, $S = \mathbb{F}_p^n$ (require $p > n$)

↳ beautiful construction based on polynomials (very useful mechanism for sharing data)

Key idea: Any d points uniquely define a degree-(d-1) polynomial over a field

- e.g. 2 points define a line, 3 points define a parabola, etc.

- given d points, can efficiently obtain entire polynomial [Lagrange interpolation]



- Share(m): Choose $y_1, \dots, y_{t-1} \xleftarrow{R} \mathbb{F}_p$

Let $f: \mathbb{F}_p \rightarrow \mathbb{F}_p$ be the unique polynomial of degree $t-1$

$f(0) = m$ and $f(i) = y_i \forall i \in [t-1]$ [t points uniquely define the polynomial f]

Output shares $(i, f(i)) \forall i \in [n]$

Each share is just 2 field elements (independent of threshold t or # parties n)

- Reconstruct $((i_1, y_1), \dots, (i_t, y_t))$: Interpolate the unique polynomial f of degree $(t-1)$ defined by the points $(i_1, y_1), \dots, (i_t, y_t)$
Output $f(0)$

Correctness: Follows by uniqueness of the interpolating polynomial

Security: Given any subset of $(t-1)$ shares $(i_1, y_1), \dots, (i_{t-1}, y_{t-1})$, and any message $m \in \mathbb{F}_p$, there is a unique polynomial f of degree $t-1$ where

$$f(i_1) = y_1, \dots, f(i_{t-1}) = y_{t-1} \text{ and } f(0) = m$$

Thus, any $(t-1)$ shares can be consistent with secret-sharing of any message $m \Rightarrow$ information-theoretic security

Efficiency: Secret sharing and secret reconstruction correspond to polynomial evaluation and interpolation (over \mathbb{F}_p). Both of these operations can be expressed as linear functions:

- Suppose $f(x) = \alpha_0 + \alpha_1 x + \dots + \alpha_{t-1} x^{t-1}$

- Evaluation of f at points x_1, x_2, \dots, x_n can be expressed as matrix-vector product:

$$\underbrace{\begin{pmatrix} 1 & x_1 & x_1^2 & x_1^3 & \dots & x_1^{t-1} \\ 1 & x_2 & x_2^2 & x_2^3 & \dots & x_2^{t-1} \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_n & x_n^2 & x_n^3 & \dots & x_n^{t-1} \end{pmatrix}}_X \underbrace{\begin{pmatrix} \alpha_0 \\ \alpha_1 \\ \vdots \\ \alpha_n \end{pmatrix}}_\alpha = \underbrace{\begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix}}_y$$


in particular,
 $y_i = f(x_i) \forall i \in [n]$

"Vandermonde matrix"

- Polynomial evaluation corresponds to computing product $X \cdot \alpha$ and interpolation corresponds to computing the product $X^{-1} y$

- Both of these operations can be done in time $O(n \log n)$ using the Fast Fourier Transform (FFT)

Computing on secret-shared data: Another paradigm for 2PC (and MPC) - better-suited for evaluating arithmetic circuits



Alice: Chooses $r_{AB}, r_{AC} \xleftarrow{\mathbb{F}_p}$ and sends r_{AB} to Bob, r_{AC} to Charlie

↳ Observation: $(x_A - r_{AB} - r_{AC}, r_{AB}, r_{AC})$ is additive secret sharing of Alice's input x_A

We will write $[x_A]$ to denote additive secret sharing of x_A

Computing on shares: Given shares of x_A and x_B ,

$$[x_A + x_B] = [x_A] + [x_B] \quad (\text{component-wise addition})$$

Specifically if $[x_A] = (x_{A,1}, x_{A,2}, x_{A,3})$ where $x_{A,1} + x_{A,2} + x_{A,3} = x_A \in \mathbb{F}_p$

$[x_B] = (x_{B,1}, x_{B,2}, x_{B,3})$ where $x_{B,1} + x_{B,2} + x_{B,3} = x_B \in \mathbb{F}_p$

then $[x_A + x_B] = (x_{A,1} + x_{B,1}, x_{A,2} + x_{B,2}, x_{A,3} + x_{B,3})$ and $(x_{A,1} + x_{B,1}) + (x_{A,2} + x_{B,2}) + (x_{A,3} + x_{B,3}) = x_A + x_B \in \mathbb{F}_p$

- More generally:
- Share addition: $[x_A + x_B] = [x_A] + [x_B]$
 - Scalar multiplication: $[k x_A] = k \cdot [x_A]$
 - Addition by constant: $[x_A + k] = (x_{A,1} + k, x_{A,2}, x_{A,3})$

Question: How to multiply shared values?

Amazing insight due to Beaver: Suppose parties have a secret-sharing of a random product: $[a], [b], [c]$ where $c = ab \in \mathbb{F}_p$

Then, given $[x]$ and $[y]$, we proceed as follows:

- Each party computes $[x-a]$ and publishes their share of $x-a$
- Each party computes $[y-b]$ and publishes their share of $y-b$
- All of the parties compute non-interactively:

$$[z] = [c] + [x](y-b) + [y](x-a) - (x-a)(y-b)$$

Claim: $z = xy$. Follows by following calculation:

$$\begin{aligned} z &= c + x(y-b) + y(x-a) - (x-a)(y-b) \\ &= \cancel{ab} + xy - \cancel{bx} + \cancel{xy} - \cancel{ay} - \cancel{xy} + \cancel{bx} + \cancel{ay} - \cancel{ab} \\ &= xy \end{aligned}$$

Observe: Parties only see $x-a$ and $y-b$ in this protocol. Since a, b are uniformly random and unknown to the parties, $x-a$ and $y-b$ is a one-time pad encryption of x and y . Resulting protocol provides information-theoretic privacy for parties' inputs.

Assuming we have access to Beaver multiplication triples, we can evaluate any arithmetic circuit as follows (among n -parties):

- Every party secret shares their input with every other party
- For each addition gate in the circuit, parties locally compute on their shares
- For each multiplication gate in the circuit, parties run Beaver's multiplication protocol (using different triple each time!)
- Every party publishes share of the output; parties run share reduction to obtain output.

What goes wrong if the same triple is used twice?

Comparison with Yao:

	Secret Sharing	Yao
Type of computation	Arithmetic circuits (\mathbb{F}_p)	Boolean circuits
Number of parties	Arbitrary (n)	2
Round complexity	Depth of circuit	2
Communication	$\sim 2n \log p$ bits per multiplication gate	~ 256 bits per AND gate*
Security	Information-theoretic (with Beaver triples)	Computational

* Can be improved further!

* Leverages several optimizations (half-gates + free XOR)

Aside: preprocessing is also possible for OT (OT correlations): gives fast, information-theoretic OT in the online phase

Question: Where do Beaver triples come from?

preprocessing is input-independent!

Typically, generated in an offline "preprocessing" phase. Many techniques:

- Trusted dealer (e.g., Intel SGX) can generate them
- Using oblivious transfer (OT) - accelerate using OT extensions (HW2, Problem 4)
- Using somewhat homomorphic encryption (discussed later in the course)

Implication: OT is "complete" for MPC: any n -party functionality can be computed securely assuming existence of OT.

Question: What if parties are malicious? [So far, everything is only semi-honest secure]

[Goldreich-Micali-Wigderson]

The GMW compiler transforms any protocol with semi-honest security to one with malicious security

High-level idea: Parties include a zero-knowledge proof that each step of the protocol was performed according to the protocol specification.

Corollary: Anything that can be computed with a trusted party can be computed without!

[Ben-Or - Goldwasser - Wigderson]

- For n -parties, if we have fewer than $n/3$ corrupted parties, this is even possible information-theoretically!
- With cryptographic assumptions (i.e. OT), we can support $n-1$ corrupted parties (with some caveats...)